

Project1-Blank

April 13, 2021

0.1 Introduction

Welcome to **CS188 - Data Science Fundamentals!** This course is designed to equip you with the tools and experiences necessary to start you off on a life-long exploration of datascience. We do not assume a prerequisite knowledge or experience in order to take the course.

For this first project we will introduce you to the end-to-end process of doing a datascience project. Our goals for this project are to:

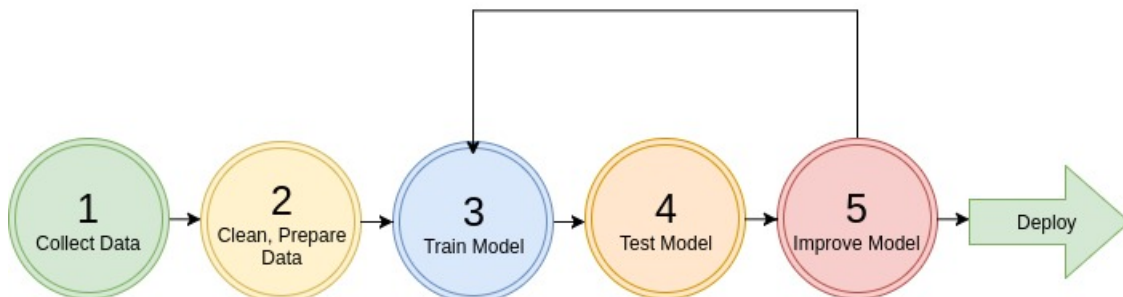
1. Familiarize you with the development environment for doing datascience
2. Get you comfortable with the python coding required to do datascience
3. Provide you with an sample end-to-end project to help you visualize the steps needed to complete a project on your own
4. Ask you to recreate a similar project on a separate dataset

In this project you will work through an example project end to end. Many of the concepts you will encounter will be unclear to you. That is OK! The course is designed to teach you these concepts in further detail. For now our focus is simply on having you replicate the code successfully and seeing a project through from start to finish.

Here are the main steps:

1. Get the data
2. Visualize the data for insights
3. Preprocess the data for your machine learning algorithm
4. Select a model and train
5. Does it meet the requirements? Fine tune the model

Steps to Machine Learning



0.2 Working with Real Data

It is best to experiment with real-data as opposed to artificial datasets.

There are many different open datasets depending on the type of problems you might be interested in!

Here are a few data repositories you could check out: - [UCI Datasets](#) - [Kaggle Datasets](#) - [AWS Datasets](#)

0.3 Submission Instructions

When you have completed this assignment please save the notebook as a PDF file and submit the assignment via Gradescope

1 Example Datascience Exercise

Below we will run through an California Housing example collected from the 1990's.

1.1 Setup

```
[2]: import sys
assert sys.version_info >= (3, 5) # python>=3.5
import sklearn
assert sklearn.__version__ >= "0.20" # sklearn >= 0.20

import numpy as np #numerical package in python
import os
%matplotlib inline
import matplotlib.pyplot as plt #plotting package

# to make this notebook's output identical at every run
np.random.seed(42)

#matplotlib magic for inline figures
%matplotlib inline
import matplotlib # plotting library
import matplotlib.pyplot as plt

# Where to save the figures
ROOT_DIR = "."
IMAGES_PATH = os.path.join(ROOT_DIR, "images")
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_name, tight_layout=True, fig_extension="png", resolution=300):
    """
    plt.savefig wrapper. refer to
    https://matplotlib.org/3.1.1/api/_as_gen/matplotlib.pyplot.savefig.html
    """
```

```

    Args:
        fig_name (str): name of the figure
        tight_layout (bool): adjust subplot to fit in the figure area
        fig_extension (str): file format to save the figure in
        resolution (int): figure resolution
    """
    path = os.path.join(IMAGES_PATH, fig_name + "." + fig_extension)
    print("Saving figure", fig_name)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)

```

```

[3]: import os
import tarfile
import urllib
DATASET_PATH = os.path.join("datasets", "housing")

```

1.2 Step 1. Getting the data

1.2.1 Intro to Data Exploration Using Pandas

In this section we will load the dataset, and visualize different features using different types of plots.

Packages we will use: - **Pandas**: is a fast, flexible and expressive data structure widely used for tabular and multidimensional datasets. - **Matplotlib**: is a 2d python plotting library which you can use to create quality figures (you can plot almost anything if you're willing to code it out!) - other plotting libraries: [seaborn](#), [ggplot2](#)

```

[4]: import pandas as pd

def load_housing_data(housing_path):
    """
        loads housing.csv dataset stored

    Args:
        housing_path (str): path to folder containing housing dataset

    Returns:
        pd.DataFrame
    """
    csv_path = os.path.join(housing_path, "housing.csv")
    return pd.read_csv(csv_path)

```

```

[5]: pd.DataFrame

```

```

[5]: pandas.core.frame.DataFrame

```

```
[6]: housing = load_housing_data(DATASET_PATH) # we load the pandas dataframe
housing.head() # show the first few elements of the dataframe
              # typically this is the first thing you do
              # to see how the dataframe looks like
```

```
[6]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

A dataset may have different types of features - real valued - Discrete (integers) - categorical (strings)

The two categorical features are essentially the same as you can always map a categorical string/character to an integer.

In the dataset example, all our features are real valued floats, except ocean proximity which is categorical.

```
[7]: # to see a concise summary of data types, null values, and counts
      # use the info() method on the dataframe
housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households              20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
[8]: # you can access individual columns similarly
# to accessing elements in a python dict
housing["ocean_proximity"].head() # added head() to avoid printing many columns.
↳.
```

```
[8]: 0    NEAR BAY
1    NEAR BAY
2    NEAR BAY
3    NEAR BAY
4    NEAR BAY
Name: ocean_proximity, dtype: object
```

```
[9]: # to access a particular row we can use iloc
housing.iloc[1]
```

```
[9]: longitude          -122.22
latitude              37.86
housing_median_age    21
total_rooms           7099
total_bedrooms        1106
population            2401
households            1138
median_income         8.3014
median_house_value    358500
ocean_proximity       NEAR BAY
Name: 1, dtype: object
```

```
[10]: # one other function that might be useful is
# value_counts(), which counts the number of occurrences
# for categorical features
housing["ocean_proximity"].value_counts()
```

```
[10]: <1H OCEAN    9136
INLAND        6551
NEAR OCEAN    2658
NEAR BAY      2290
ISLAND         5
Name: ocean_proximity, dtype: int64
```

```
[11]: # The describe function compiles your typical statistics for each
# column
housing.describe()
```

```
[11]:
```

	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	

min	-124.350000	32.540000	1.000000	2.000000
25%	-121.800000	33.930000	18.000000	1447.750000
50%	-118.490000	34.260000	29.000000	2127.000000
75%	-118.010000	37.710000	37.000000	3148.000000
max	-114.310000	41.950000	52.000000	39320.000000

	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	
mean	537.870553	1425.476744	499.539680	3.870671	
std	421.385070	1132.462122	382.329753	1.899822	
min	1.000000	3.000000	1.000000	0.499900	
25%	296.000000	787.000000	280.000000	2.563400	
50%	435.000000	1166.000000	409.000000	3.534800	
75%	647.000000	1725.000000	605.000000	4.743250	
max	6445.000000	35682.000000	6082.000000	15.000100	

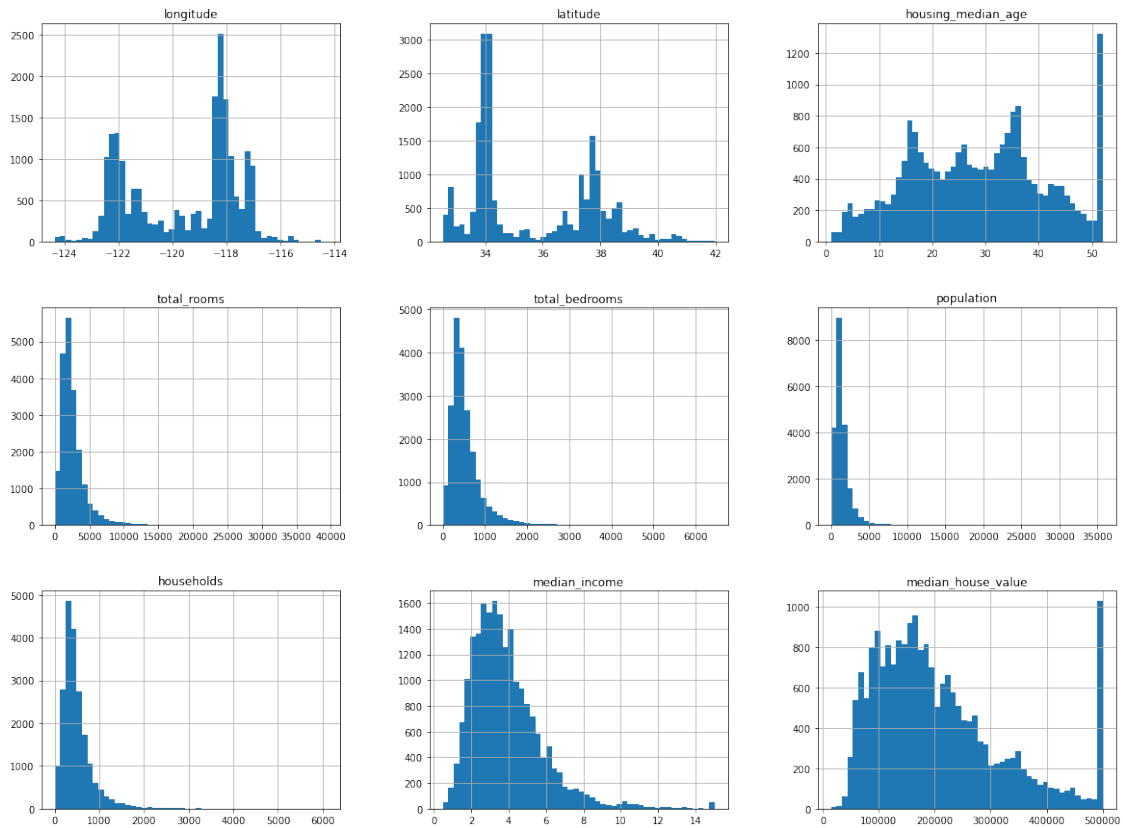
	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

If you want to learn about different ways of accessing elements or other functions it's useful to check out the getting started section [here](#)

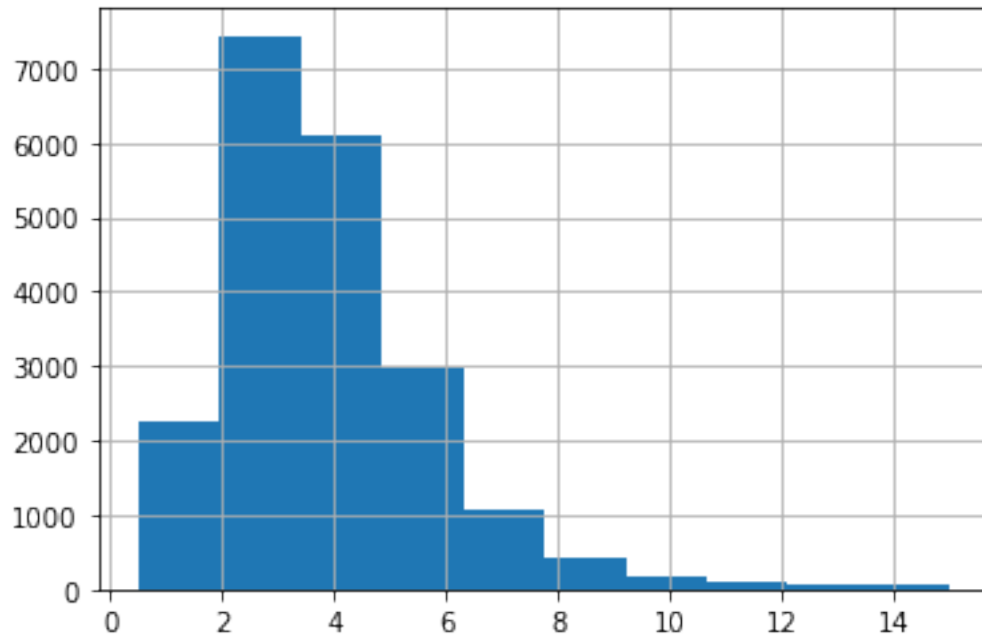
1.3 Step 2. Visualizing the data

1.3.1 Let's start visualizing the dataset

```
[12]: # We can draw a histogram for each of the dataframes features
      # using the hist function
housing.hist(bins=50, figsize=(20,15))
      # save_fig("attribute_histogram_plots")
plt.show() # pandas internally uses matplotlib, and to display all the figures
           # the show() function must be called
```



```
[13]: # if you want to have a histogram on an individual feature:
housing["median_income"].hist() # default is 10 bins
plt.show()
```



We can convert a floating point feature to a categorical feature by binning or by defining a set of intervals.

For example, to bin the households based on median_income we can use the pd.cut function

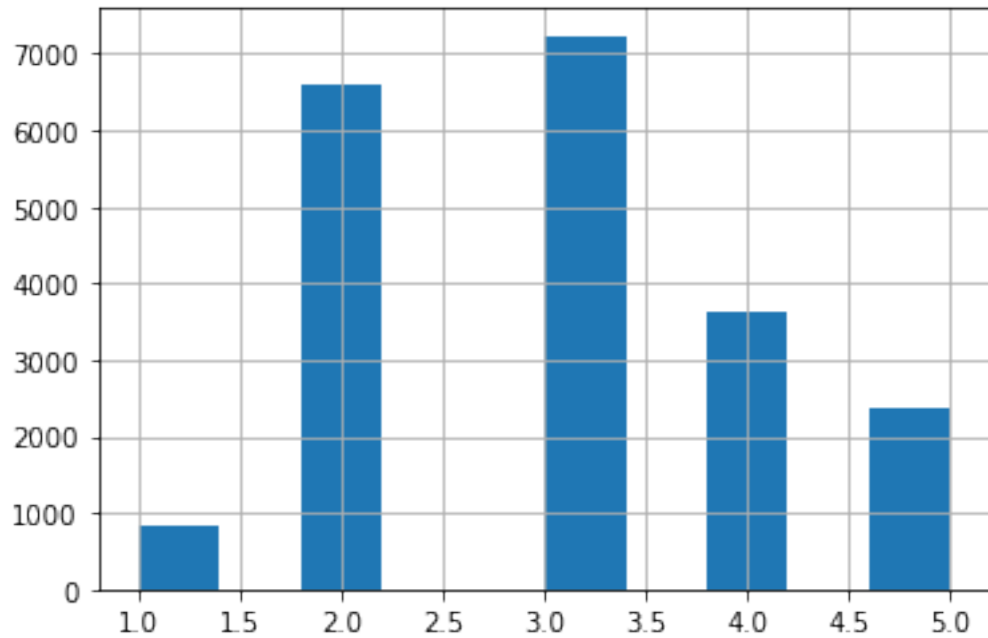
```
[14]: # assign each bin a categorical value [1, 2, 3, 4, 5] in this case.
housing["income_cat"] = pd.cut(housing["median_income"],
                               bins=[0., 1.5, 3.0, 4.5, 6., np.inf],
                               labels=[1, 2, 3, 4, 5])

housing["income_cat"].value_counts()
```

```
[14]: 3    7236
      2    6581
      4    3639
      5    2362
      1     822
      Name: income_cat, dtype: int64
```

```
[15]: housing["income_cat"].hist()
```

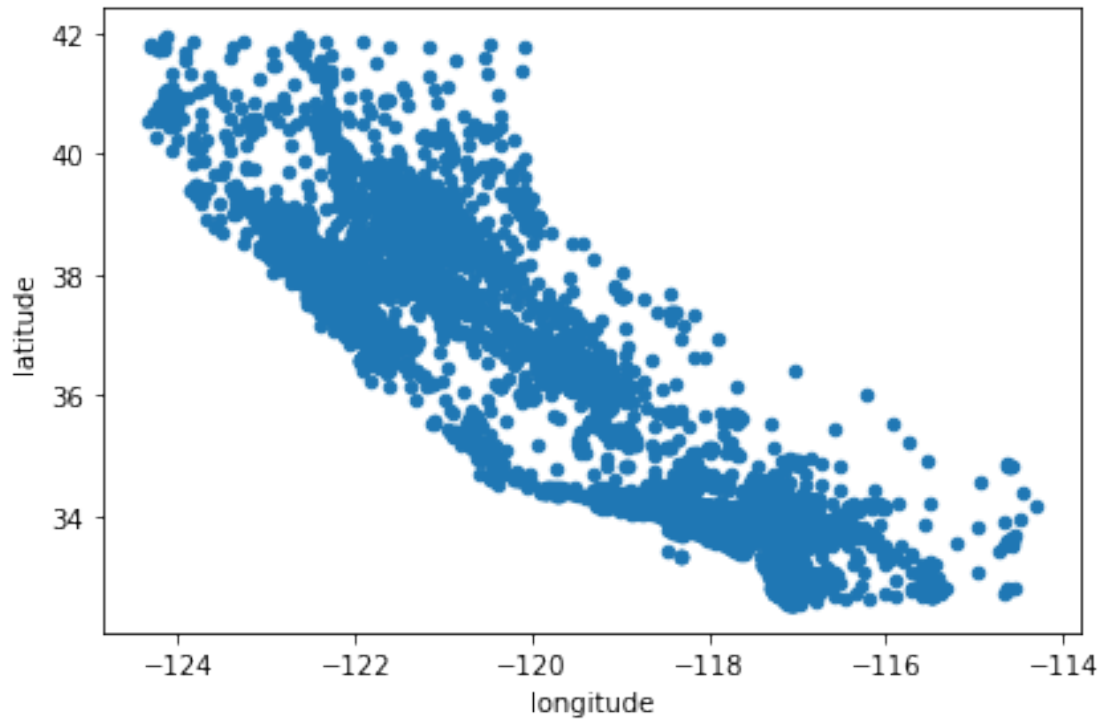
```
[15]: <AxesSubplot:>
```

Next let's visualize the household incomes based on latitude & longitude coordinates

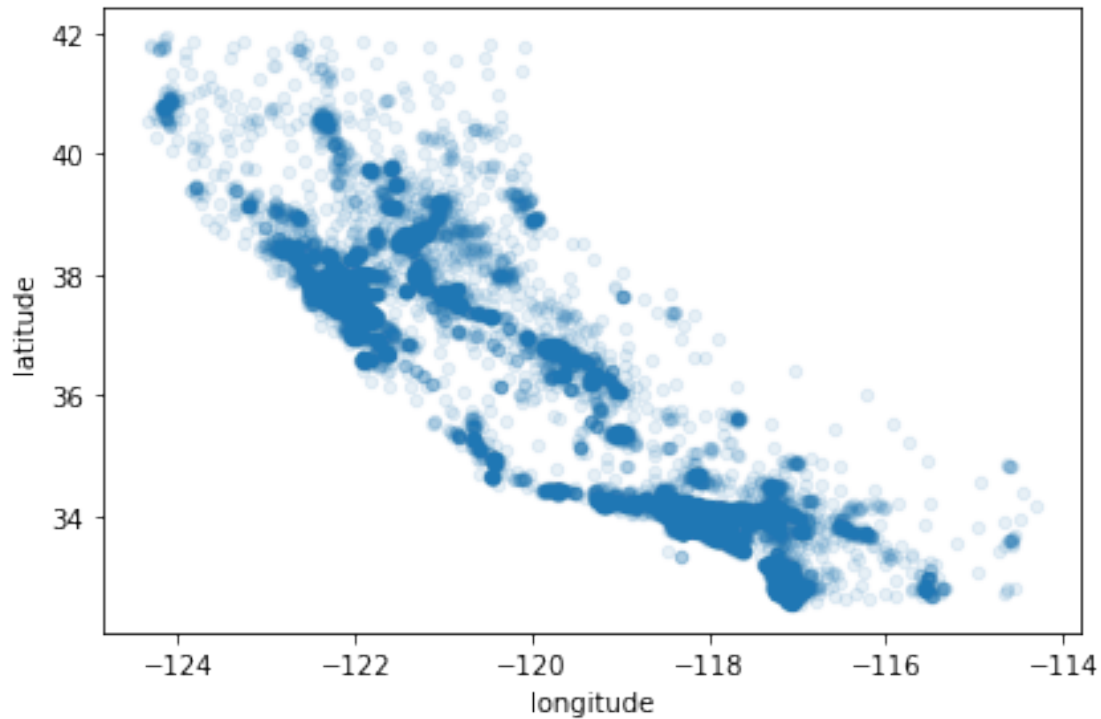
```
[16]: ## here's a not so interesting way of plotting it  
housing.plot(kind="scatter", x="longitude", y="latitude")  
save_fig("bad_visualization_plot")
```

Saving figure bad_visualization_plot



```
[17]: # we can make it look a bit nicer by using the alpha parameter,  
# it simply plots less dense areas lighter.  
housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.1)  
save_fig("better_visualization_plot")
```

Saving figure better_visualization_plot



```
[18]: # A more interesting plot is to color code (heatmap) the dots
# based on income. The code below achieves this

# load an image of california
images_path = os.path.join('.', "images")
os.makedirs(images_path, exist_ok=True)
filename = "california.png"

import matplotlib.image as mpimg
california_img=mpimg.imread(os.path.join(images_path, filename))
ax = housing.plot(kind="scatter", x="longitude", y="latitude", figsize=(10,7),
                  s=housing['population']/100, label="Population",
                  c="median_house_value", cmap=plt.get_cmap("jet"),
                  colorbar=False, alpha=0.4,
                  )
# note above how if we remove colorbar=False above, a duplicate colorbar will
# appear

# overlay the califronia map on the plotted scatter plot
# note: plt.imshow still refers to the most recent figure
# that hasn't been plotted yet.
plt.imshow(california_img, extent=[-124.55, -113.80, 32.45, 42.05], alpha=0.5,
           cmap=plt.get_cmap("jet"))
```

```

plt.ylabel("Latitude", fontsize=14)
plt.xlabel("Longitude", fontsize=14)

# setting up heatmap colors based on median_house_value feature
prices = housing["median_house_value"]
tick_values = np.linspace(prices.min(), prices.max(), 11)
print(tick_values)
cb = plt.colorbar() # add a colorbar to plot
# %d is a formatter for integers; k is to represent "thousand" in the scale
cb.ax.set_yticklabels(["$%dk"%(round(v/1000)) for v in tick_values],
    ↳fontsize=14)
cb.set_label('Median House Value', fontsize=16)

# Why are there only 7 ticks in the colorbar below but our linspace function
↳above divided it into 11 evenly-spaced tick values???

plt.legend(fontsize=16)
save_fig("california_housing_prices_plot")
plt.show()

```

```

[ 14999.    63499.2 111999.4 160499.6 208999.8 257500.   306000.2 354500.4
 403000.6 451500.8 500001. ]

```

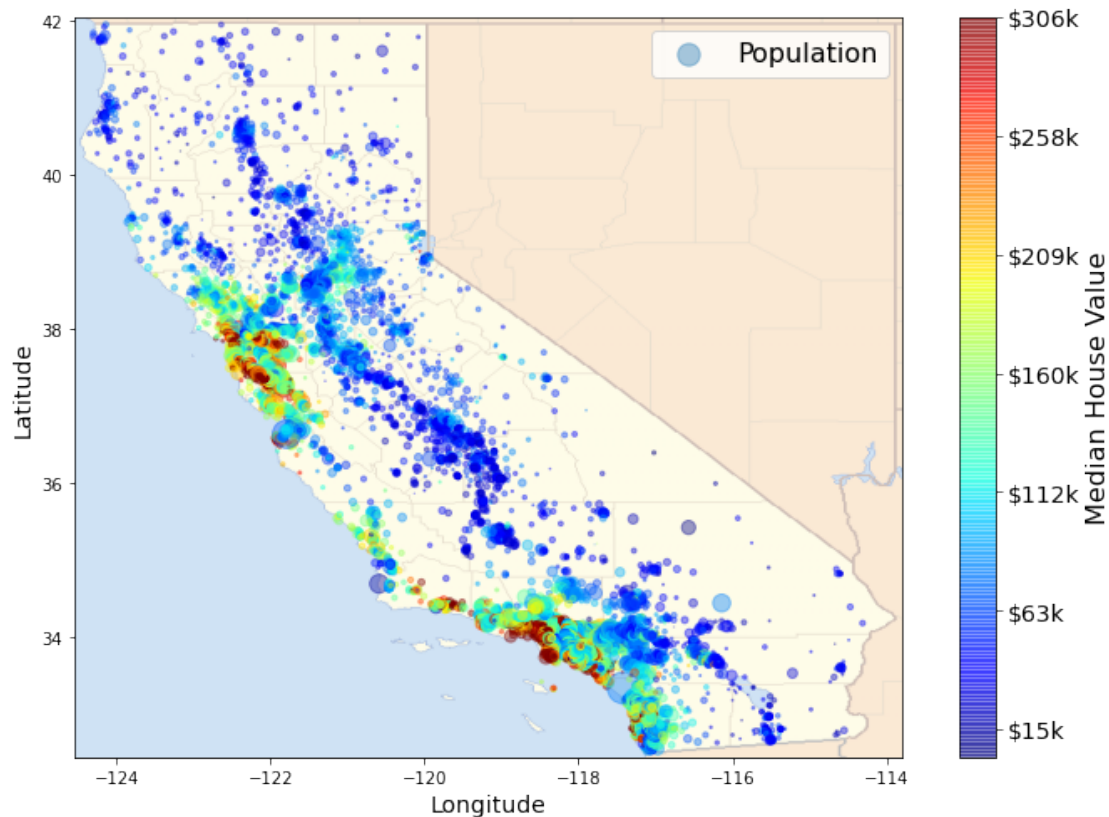
Saving figure california_housing_prices_plot

<ipython-input-18-f361eec34593>:32: UserWarning: FixedFormatter should only be used together with FixedLocator

```

    cb.ax.set_yticklabels(["$%dk"%(round(v/1000)) for v in tick_values],
    fontsize=14)

```



Not surprisingly, we can see that the most expensive houses are concentrated around the San Francisco/Los Angeles areas.

Up until now we have only visualized feature histograms and basic statistics.

When developing machine learning models the predictiveness of a feature for a particular target of interest is what's important.

It may be that only a few features are useful for the target at hand, or features may need to be augmented by applying certain transformations.

None the less we can explore this using correlation matrices. If you need to brush up on correlation take a look [here](#).

```
[19]: corr_matrix = housing.corr() # compute the correlation matrix
```

```
[20]: # for example if the target is "median_house_value", most correlated features
      ↪ can be sorted
      # which happens to be "median_income". This also intuitively makes sense.
      corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
[20]: median_house_value    1.000000
      median_income        0.688075
```

```

total_rooms      0.134153
housing_median_age 0.105623
households       0.065843
total_bedrooms   0.049686
population       -0.024650
longitude        -0.045967
latitude         -0.144160
Name: median_house_value, dtype: float64

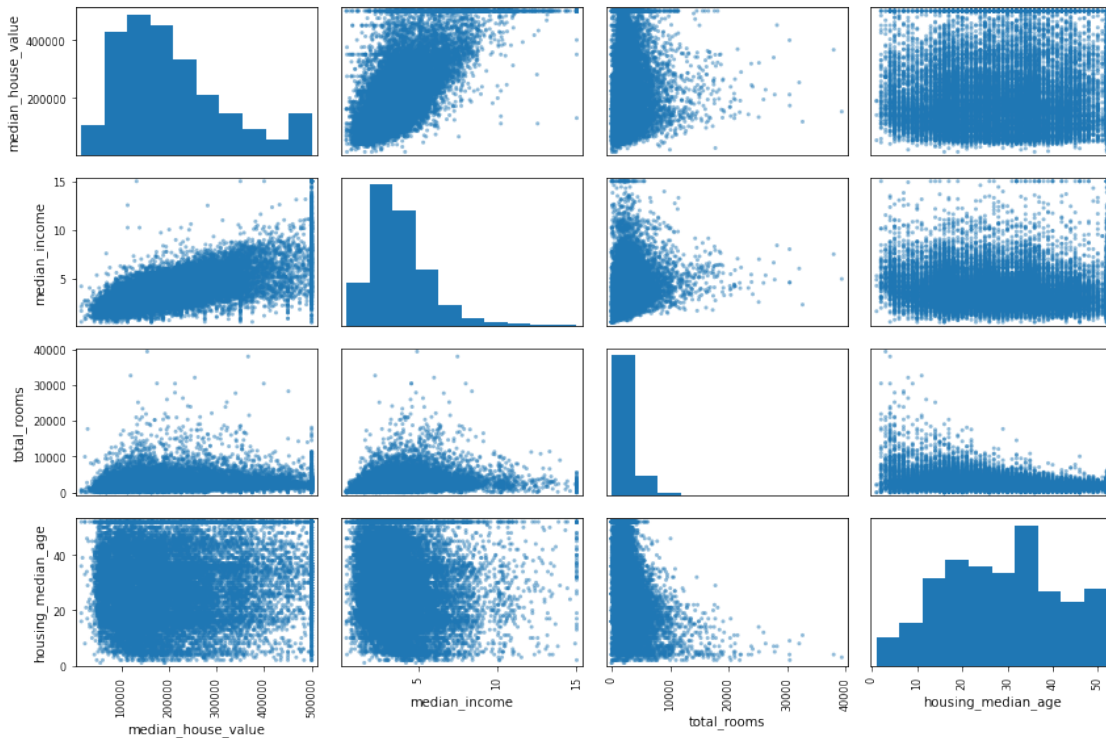
```

```

[21]: # the correlation matrix for different attributes/features can also be plotted
# some features may show a positive correlation/negative correlation or
# it may turn out to be completely random!
from pandas.plotting import scatter_matrix
attributes = ["median_house_value", "median_income", "total_rooms",
             "housing_median_age"]
scatter_matrix(housing[attributes], figsize=(12, 8))
save_fig("scatter_matrix_plot")

```

Saving figure scatter_matrix_plot



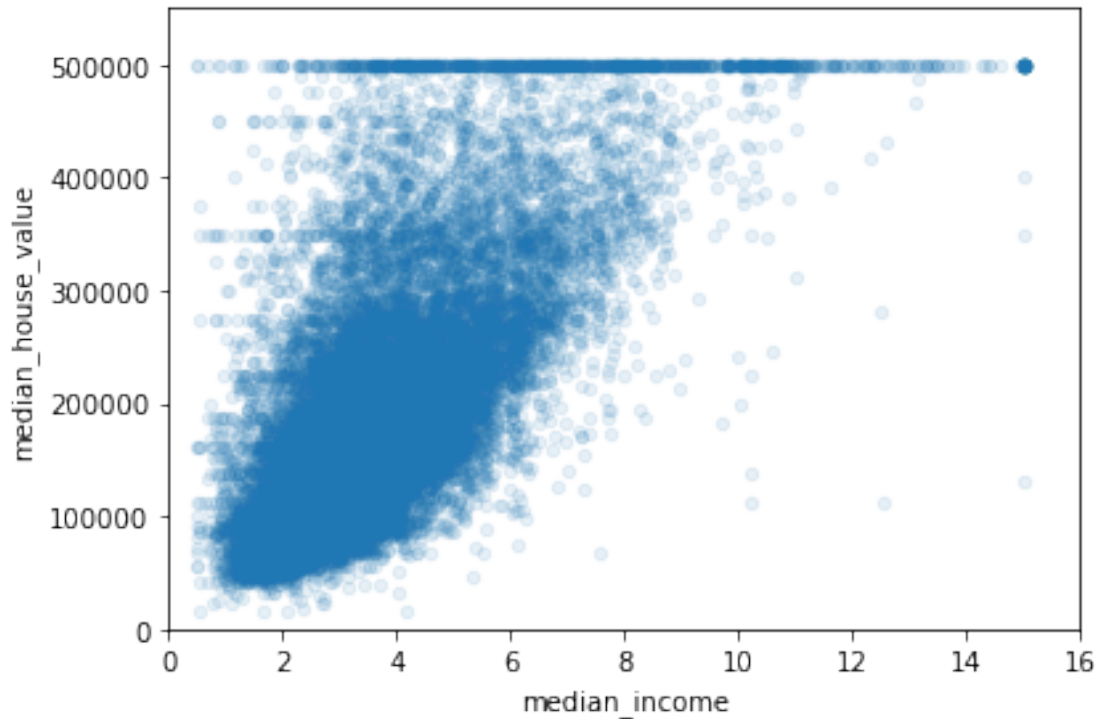
```

[22]: # median income vs median house value plot plot 2 in the first row of top figure
housing.plot(kind="scatter", x="median_income", y="median_house_value",
             alpha=0.1)

```

```
plt.axis([0, 16, 0, 550000])
save_fig("income_vs_house_value_scatterplot")
```

Saving figure income_vs_house_value_scatterplot



1.3.2 Augmenting Features

New features can be created by combining different columns from our data set.

- $\text{rooms_per_household} = \text{total_rooms} / \text{households}$
- $\text{bedrooms_per_room} = \text{total_bedrooms} / \text{total_rooms}$
- etc.

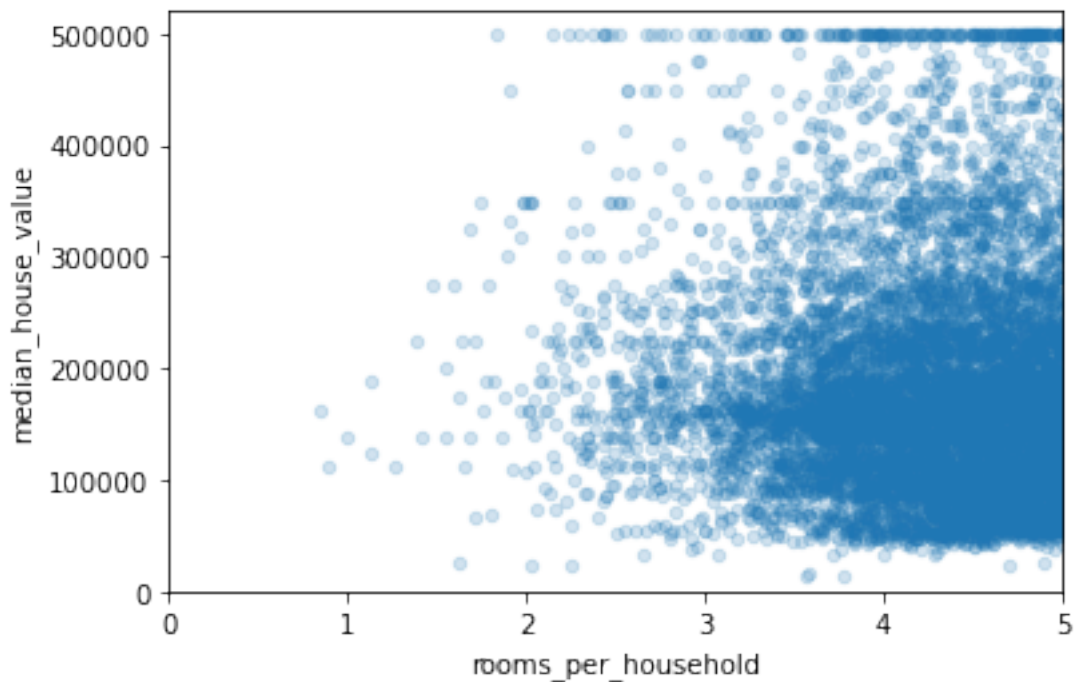
```
[23]: housing["rooms_per_household"] = housing["total_rooms"]/housing["households"]
housing["bedrooms_per_room"] = housing["total_bedrooms"]/housing["total_rooms"]
housing["population_per_household"]=housing["population"]/housing["households"]
```

```
[24]: # obtain new correlations
corr_matrix = housing.corr()
corr_matrix["median_house_value"].sort_values(ascending=False)
```

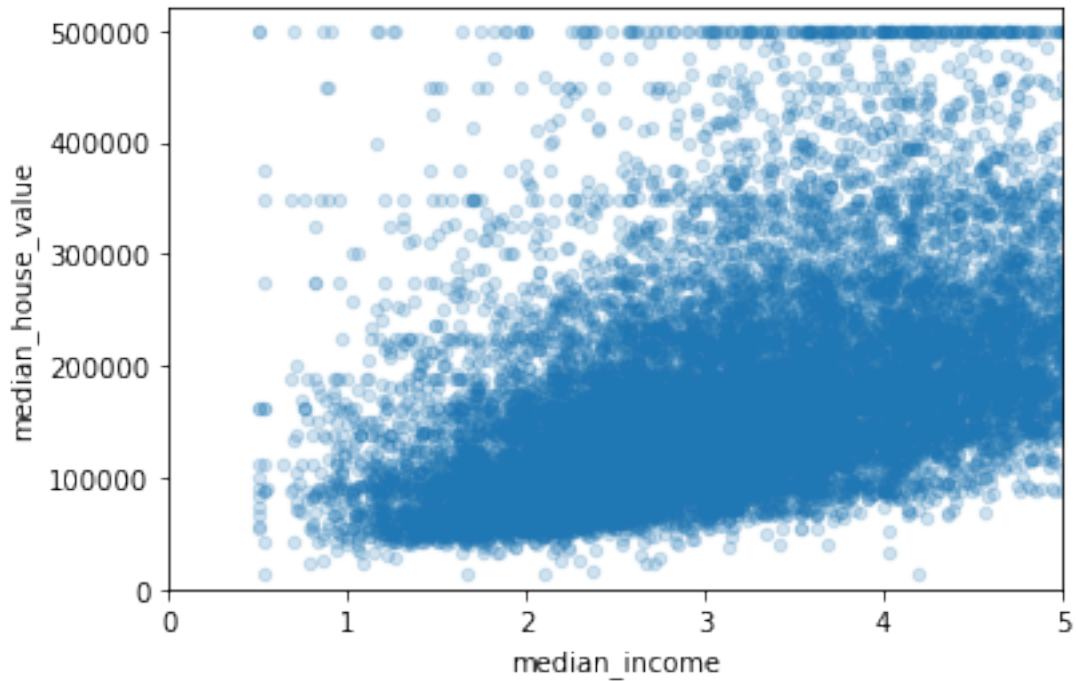
```
[24]: median_house_value    1.000000
      median_income        0.688075
      rooms_per_household   0.151948
```

```
total_rooms          0.134153
housing_median_age    0.105623
households            0.065843
total_bedrooms        0.049686
population_per_household -0.023737
population            -0.024650
longitude             -0.045967
latitude              -0.144160
bedrooms_per_room     -0.255880
Name: median_house_value, dtype: float64
```

```
[25]: housing.plot(kind="scatter", x="rooms_per_household", y="median_house_value",
                    alpha=0.2)
plt.axis([0, 5, 0, 520000])
plt.show()
```



```
[26]: housing.plot(kind="scatter", x="median_income", y="median_house_value",
                    alpha=0.2)
plt.axis([0, 5, 0, 520000])
plt.show()
```

```
[27]: housing.describe()
```

```
[27]:
```

	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	
mean	537.870553	1425.476744	499.539680	3.870671	
std	421.385070	1132.462122	382.329753	1.899822	
min	1.000000	3.000000	1.000000	0.499900	
25%	296.000000	787.000000	280.000000	2.563400	
50%	435.000000	1166.000000	409.000000	3.534800	
75%	647.000000	1725.000000	605.000000	4.743250	
max	6445.000000	35682.000000	6082.000000	15.000100	

	median_house_value	rooms_per_household	bedrooms_per_room	\
count	20640.000000	20640.000000	20433.000000	

mean	206855.816909	5.429000	0.213039
std	115395.615874	2.474173	0.057983
min	14999.000000	0.846154	0.100000
25%	119600.000000	4.440716	0.175427
50%	179700.000000	5.229129	0.203162
75%	264725.000000	6.052381	0.239821
max	500001.000000	141.909091	1.000000

	population_per_household
count	20640.000000
mean	3.070655
std	10.386050
min	0.692308
25%	2.429741
50%	2.818116
75%	3.282261
max	1243.333333

1.4 Step 3. Preprocess the data for your machine learning algorithm

Once we've visualized the data, and have a certain understanding of how the data looks like. It's time to clean!

Most of your time will be spent on this step, although the datasets used in this project are relatively nice and clean... in the real world it could get real dirty.

After having cleaned your dataset you're aiming for: - train set - test set

In some cases you might also have a validation set as well for tuning hyperparameters (don't worry if you're not familiar with this term yet..)

In supervised learning setting your train set and test set should contain (**feature**, **target**) tuples. - **feature**: is the input to your model - **target**: is the ground truth label - when target is categorical the task is a classification task - when target is floating point the task is a regression task - I don't really understand this. Why is it a floating point, and what qualifies as a regression task for supervised learning???

We will make use of [scikit-learn](#) python package for preprocessing.

Scikit learn is pretty well documented and if you get confused at any point simply look up the function/object!

1.4.1 Dealing With Incomplete Data

```
[28]: # have you noticed when looking at the dataframe summary certain rows
      # contained null values? we can't just leave them as nulls and expect our
      # model to handle them for us so we'll have to devise a method for dealing with
      # them...
      # I'm a little confused by what pd.any() does???
      sample_incomplete_rows = housing[housing.isnull().any(axis=1)].head()
```

```
sample_incomplete_rows
```

```
[28]:      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
290      -122.16    37.77                47.0      1256.0             NaN
341      -122.17    37.75                38.0       992.0             NaN
538      -122.28    37.78                29.0     5154.0             NaN
563      -122.24    37.75                45.0       891.0             NaN
696      -122.10    37.69                41.0       746.0             NaN

      population  households  median_income  median_house_value  \
290         570.0        218.0         4.3750         161900.0
341         732.0        259.0         1.6196         85100.0
538        3741.0       1273.0         2.5762        173400.0
563         384.0        146.0         4.9489        247100.0
696         387.0        161.0         3.9063        178400.0

      ocean_proximity  income_cat  rooms_per_household  bedrooms_per_room  \
290          NEAR BAY           3          5.761468             NaN
341          NEAR BAY           2          3.830116             NaN
538          NEAR BAY           2          4.048704             NaN
563          NEAR BAY           4          6.102740             NaN
696          NEAR BAY           3          4.633540             NaN

      population_per_household
290              2.614679
341              2.826255
538              2.938727
563              2.630137
696              2.403727
```

```
[29]: sample_incomplete_rows.dropna(subset=["total_bedrooms"])    # option 1: simply
      ↪ drop rows that have null values
```

```
[29]: Empty DataFrame
Columns: [longitude, latitude, housing_median_age, total_rooms, total_bedrooms,
population, households, median_income, median_house_value, ocean_proximity,
income_cat, rooms_per_household, bedrooms_per_room, population_per_household]
Index: []
```

```
[30]: sample_incomplete_rows.drop("total_bedrooms", axis=1)      # option 2: drop
      ↪ the complete feature
```

```
[30]:      longitude  latitude  housing_median_age  total_rooms  population  \
290      -122.16    37.77                47.0      1256.0         570.0
341      -122.17    37.75                38.0       992.0         732.0
538      -122.28    37.78                29.0     5154.0        3741.0
563      -122.24    37.75                45.0       891.0         384.0
```

696	-122.10	37.69	41.0	746.0	387.0
-----	---------	-------	------	-------	-------

	households	median_income	median_house_value	ocean_proximity	income_cat \
290	218.0	4.3750	161900.0	NEAR BAY	3
341	259.0	1.6196	85100.0	NEAR BAY	2
538	1273.0	2.5762	173400.0	NEAR BAY	2
563	146.0	4.9489	247100.0	NEAR BAY	4
696	161.0	3.9063	178400.0	NEAR BAY	3

	rooms_per_household	bedrooms_per_room	population_per_household
290	5.761468	NaN	2.614679
341	3.830116	NaN	2.826255
538	4.048704	NaN	2.938727
563	6.102740	NaN	2.630137
696	4.633540	NaN	2.403727

```
[31]: median = housing["total_bedrooms"].median()
sample_incomplete_rows["total_bedrooms"].fillna(median, inplace=True) # option 1
↳ 3: replace na values with median values
sample_incomplete_rows
```

```
[31]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms \
290	-122.16	37.77	47.0	1256.0	435.0
341	-122.17	37.75	38.0	992.0	435.0
538	-122.28	37.78	29.0	5154.0	435.0
563	-122.24	37.75	45.0	891.0	435.0
696	-122.10	37.69	41.0	746.0	435.0

	population	households	median_income	median_house_value \
290	570.0	218.0	4.3750	161900.0
341	732.0	259.0	1.6196	85100.0
538	3741.0	1273.0	2.5762	173400.0
563	384.0	146.0	4.9489	247100.0
696	387.0	161.0	3.9063	178400.0

	ocean_proximity	income_cat	rooms_per_household	bedrooms_per_room \
290	NEAR BAY	3	5.761468	NaN
341	NEAR BAY	2	3.830116	NaN
538	NEAR BAY	2	4.048704	NaN
563	NEAR BAY	4	6.102740	NaN
696	NEAR BAY	3	4.633540	NaN

	population_per_household
290	2.614679
341	2.826255
538	2.938727
563	2.630137

Could you think of another plausible imputation for this dataset? (Not graded)

1.4.2 Prepare Data

Recall we are trying to predict the median house value, our features will contain longitude, latitude, housing_median_age... and our target will be median_house_value

```
[32]: housing_features = housing.drop("median_house_value", axis=1) # drop labels for
      ↪ training set features
      # the input to the model
      ↪ should not contain the true label
      housing_labels = housing["median_house_value"].copy()
```

```
[33]: housing_features.head()
```

```
[33]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	ocean_proximity	income_cat	\
0	322.0	126.0	8.3252	NEAR BAY	5	
1	2401.0	1138.0	8.3014	NEAR BAY	5	
2	496.0	177.0	7.2574	NEAR BAY	5	
3	558.0	219.0	5.6431	NEAR BAY	4	
4	565.0	259.0	3.8462	NEAR BAY	3	

	rooms_per_household	bedrooms_per_room	population_per_household
0	6.984127	0.146591	2.555556
1	6.238137	0.155797	2.109842
2	8.288136	0.129516	2.802260
3	5.817352	0.184458	2.547945
4	6.281853	0.172096	2.181467

```
[34]: # This cell implements the complete pipeline for preparing the data
      # using sklearn's TransformerMixins
      # Earlier we mentioned different types of features: categorical, and floats.
      # In the case of floats we might want to convert them to categories.
      # On the other hand categories in which are not already represented as integers
      ↪ must be mapped to integers before
      # feeding to the model.
```

```

# Additionally, categorical values could either be represented as one-hot
↳ vectors or simple as normalized/unnormalized integers.
# Here we encode them using one hot vectors.

# DO NOT WORRY IF YOU DO NOT UNDERSTAND ALL THE STEPS OF THIS PIPELINE.
↳ CONCEPTS LIKE NORMALIZATION,
# ONE-HOT ENCODING ETC. WILL ALL BE COVERED IN DISCUSSION

from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder

from sklearn.base import BaseEstimator, TransformerMixin

imputer = SimpleImputer(strategy="median") # use median imputation for missing
↳ values
housing_num = housing_features.drop("ocean_proximity", axis=1) # remove the
↳ categorical feature
# column index
rooms_idx, bedrooms_idx, population_idx, households_idx = 3, 4, 5, 6

#
class AugmentFeatures(BaseEstimator, TransformerMixin):
    """
    implements the previous features we had defined
    housing["rooms_per_household"] = housing["total_rooms"] /
    ↳ housing["households"]
    housing["bedrooms_per_room"] = housing["total_bedrooms"] /
    ↳ housing["total_rooms"]
    housing["population_per_household"] = housing["population"] /
    ↳ housing["households"]
    """
    def __init__(self, add_bedrooms_per_room = True):
        self.add_bedrooms_per_room = add_bedrooms_per_room

    def fit(self, X, y=None):
        return self # nothing else to do

    def transform(self, X):
        rooms_per_household = X[:, rooms_idx] / X[:, households_idx]
        population_per_household = X[:, population_idx] / X[:, households_idx]
        if self.add_bedrooms_per_room:

```

```

        bedrooms_per_room = X[:, bedrooms_idx] / X[:, rooms_idx]
        return np.c_[X, rooms_per_household, population_per_household,
                     bedrooms_per_room]
    else:
        return np.c_[X, rooms_per_household, population_per_household]

attr_adder = AugmentFeatures(add_bedrooms_per_room=False)
housing_extra_attribs = attr_adder.transform(housing.values) # generate new
↳ features

# this will be a numerical pipeline
# 1. impute, 2. augment the feature set 3. normalize using StandardScaler()
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attribs_adder', AugmentFeatures()),
    ('std_scaler', StandardScaler()),
])

housing_num_tr = num_pipeline.fit_transform(housing_num)

numerical_features = list(housing_num)
categorical_features = ["ocean_proximity"]

full_pipeline = ColumnTransformer([
    ("num", num_pipeline, numerical_features),
    ("cat", OneHotEncoder(), categorical_features),
])

housing_prepared = full_pipeline.fit_transform(housing_features)

```

1.4.3 Splitting our dataset

First we need to carve out our dataset into a training and testing cohort. To do this we'll use `train_test_split`, a very elementary tool that arbitrarily splits the data into training and testing cohorts.

```

[35]: from sklearn.model_selection import train_test_split
data_target = housing['median_house_value']
train, test, target, target_test = train_test_split(housing_prepared,
↳ data_target, test_size=0.3, random_state=0)

```

1.4.4 Select a model and train

Once we have prepared the dataset it's time to choose a model.

As our task is to predict the `median_house_value` (a floating value), regression is well suited for this.

```
[36]: from sklearn.linear_model import LinearRegression
```

```
lin_reg = LinearRegression()  
lin_reg.fit(train, target)
```

```
# let's try the full preprocessing pipeline on a few training instances  
data = test  
labels = target_test
```

```
print("Predictions:", lin_reg.predict(data)[:5])  
print("Actual labels:", list(labels)[:5])
```

```
Predictions: [207828.06448011 281099.80175494 176021.36890539 93643.46744928  
304674.47047758]
```

```
Actual labels: [136900.0, 241300.0, 200700.0, 72500.0, 460000.0]
```

```
[37]: from sklearn.metrics import mean_squared_error
```

```
preds = lin_reg.predict(test)  
mse = mean_squared_error(target_test, preds)  
rmse = np.sqrt(mse)  
rmse
```

```
[37]: 67879.86844243006
```

2 TODO: Applying the end-end ML steps to a different dataset.

We will apply what we've learnt to another dataset (airbnb dataset). We will predict airbnb price based on other features.

3 [35 pts] Visualizing Data

3.0.1 [5 pts] Load the data + statistics

- load the dataset
- display the first few rows of the data

```
[110]: def load_airbnb_data(airbnb_path):
```

```
    '''
```

```
        loads AB_NYC_2019.csv dataset stored
```

```
    Args:
```

```
        airbnb_path (str): path to folder containing airbnb dataset
```

```
    Returns:
```

```
        pd.DataFrame
```

```
    '''
```



```
csv_path = os.path.join(airbnb_path, "AB_NYC_2019.csv")
return pd.read_csv(csv_path)
```

```
DATASET_PATH = os.path.join("datasets", "airbnb")
airbnb = load_airbnb_data(DATASET_PATH)
airbnb.head()
```

```
[110]:
```

	id	name	host_id	\
0	2539	Clean & quiet apt home by the park	2787	
1	2595	Skylit Midtown Castle	2845	
2	3647	THE VILLAGE OF HARLEM...NEW YORK !	4632	
3	3831	Cozy Entire Floor of Brownstone	4869	
4	5022	Entire Apt: Spacious Studio/Loft by central park	7192	

	host_name	neighbourhood_group	neighbourhood	latitude	longitude	\
0	John	Brooklyn	Kensington	40.64749	-73.97237	
1	Jennifer	Manhattan	Midtown	40.75362	-73.98377	
2	Elisabeth	Manhattan	Harlem	40.80902	-73.94190	
3	LisaRoxanne	Brooklyn	Clinton Hill	40.68514	-73.95976	
4	Laura	Manhattan	East Harlem	40.79851	-73.94399	

	room_type	price	minimum_nights	number_of_reviews	last_review	\
0	Private room	149	1	9	2018-10-19	
1	Entire home/apt	225	1	45	2019-05-21	
2	Private room	150	3	0	NaN	
3	Entire home/apt	89	1	270	2019-07-05	
4	Entire home/apt	80	10	9	2018-11-19	

	reviews_per_month	calculated_host_listings_count	availability_365
0	0.21	6	365
1	0.38	2	355
2	NaN	1	365
3	4.64	1	194
4	0.10	1	0

- pull up info on the data type for each of the data fields. Will any of these be problematic feeding into your model (you may need to do a little research on this)? Discuss:

```
[111]: airbnb.info()
# As we can see below, a few fields like name, host_name, last_review, and
# reviews_per_month have some null values that we have to deal with.
# Some of the fields like neighborhood and room_type are categorical.
# What is the difference between neighborhood_group and neighborhood???
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
#   ...
```

```

---  -----
0    id                48895 non-null  int64
1    name              48879 non-null  object
2    host_id           48895 non-null  int64
3    host_name         48874 non-null  object
4    neighbourhood_group 48895 non-null  object
5    neighbourhood     48895 non-null  object
6    latitude          48895 non-null  float64
7    longitude         48895 non-null  float64
8    room_type         48895 non-null  object
9    price             48895 non-null  int64
10   minimum_nights    48895 non-null  int64
11   number_of_reviews 48895 non-null  int64
12   last_review       38843 non-null  object
13   reviews_per_month 38843 non-null  float64
14   calculated_host_listings_count 48895 non-null  int64
15   availability_365   48895 non-null  int64
dtypes: float64(3), int64(7), object(6)
memory usage: 6.0+ MB

```

[Response here]

- drop the following columns: name, host_id, host_name, and last_review
- display a summary of the statistics of the loaded data

```

[112]: airbnb.drop("name", axis=1, inplace=True)
airbnb.drop("host_id", axis=1, inplace=True)
airbnb.drop("host_name", axis=1, inplace=True)
airbnb.drop("last_review", axis=1, inplace=True)
airbnb.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  ---
0    id                48895 non-null  int64
1    neighbourhood_group 48895 non-null  object
2    neighbourhood     48895 non-null  object
3    latitude          48895 non-null  float64
4    longitude         48895 non-null  float64
5    room_type         48895 non-null  object
6    price             48895 non-null  int64
7    minimum_nights    48895 non-null  int64
8    number_of_reviews 48895 non-null  int64
9    reviews_per_month 38843 non-null  float64
10   calculated_host_listings_count 48895 non-null  int64
11   availability_365   48895 non-null  int64
dtypes: float64(3), int64(6), object(3)

```

memory usage: 4.5+ MB

```
[113]: airbnb.describe()
```

```
[113]:
```

	id	latitude	longitude	price	minimum_nights \
count	4.889500e+04	48895.000000	48895.000000	48895.000000	48895.000000
mean	1.901714e+07	40.728949	-73.952170	152.720687	7.029962
std	1.098311e+07	0.054530	0.046157	240.154170	20.510550
min	2.539000e+03	40.499790	-74.244420	0.000000	1.000000
25%	9.471945e+06	40.690100	-73.983070	69.000000	1.000000
50%	1.967728e+07	40.723070	-73.955680	106.000000	3.000000
75%	2.915218e+07	40.763115	-73.936275	175.000000	5.000000
max	3.648724e+07	40.913060	-73.712990	10000.000000	1250.000000

	number_of_reviews	reviews_per_month	calculated_host_listings_count \
count	48895.000000	38843.000000	48895.000000
mean	23.274466	1.373221	7.143982
std	44.550582	1.680442	32.952519
min	0.000000	0.010000	1.000000
25%	1.000000	0.190000	1.000000
50%	5.000000	0.720000	1.000000
75%	24.000000	2.020000	2.000000
max	629.000000	58.500000	327.000000

	availability_365
count	48895.000000
mean	112.781327
std	131.622289
min	0.000000
25%	0.000000
50%	45.000000
75%	227.000000
max	365.000000

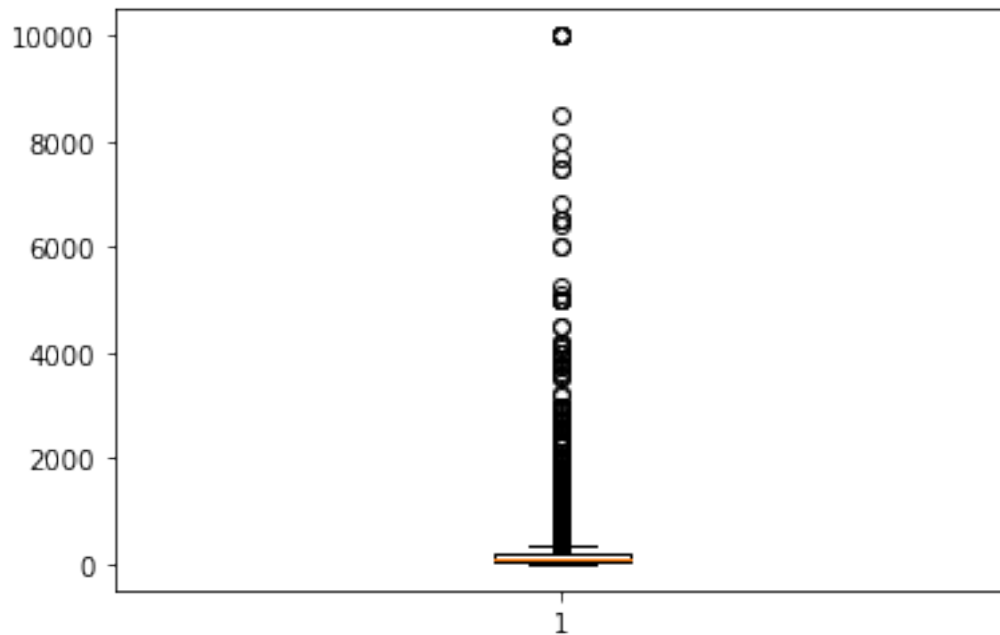
3.0.2 [5 pts] Boxplot 3 features of your choice

- plot boxplots for 3 features of your choice

```
[114]: # columns = [airbnb["price"], airbnb["minimum_nights"],  
↪ airbnb["availability_365"]]  
fig, ax = plt.subplots()  
ax.boxplot(airbnb["price"])
```

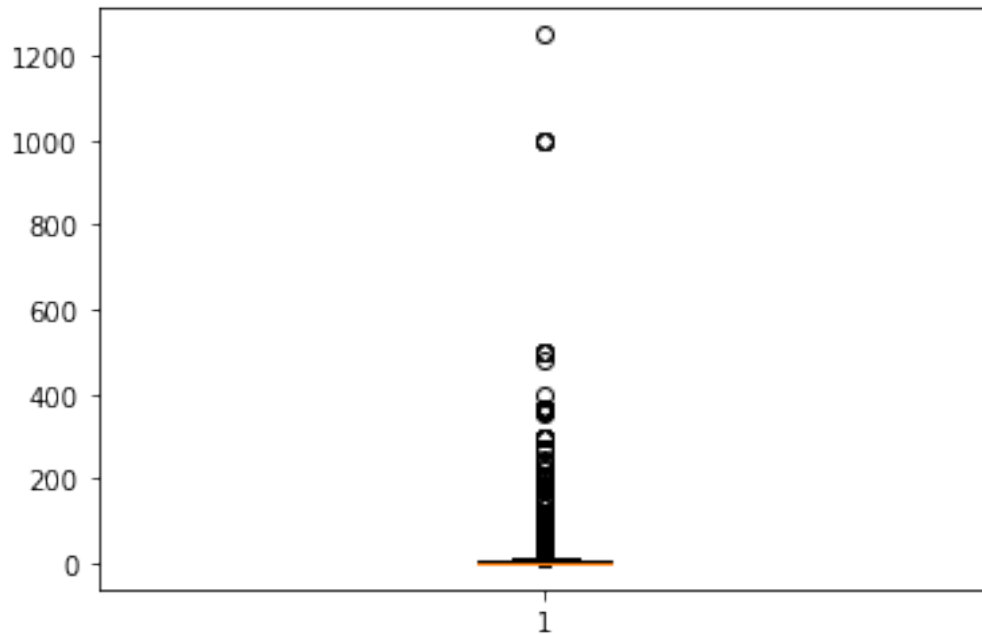
```
[114]: {'whiskers': [<matplotlib.lines.Line2D at 0x28d76d84700>,  
<matplotlib.lines.Line2D at 0x28d76d84460>],  
'caps': [<matplotlib.lines.Line2D at 0x28d72bbbf0>,  
<matplotlib.lines.Line2D at 0x28d72bbb340>],  
'boxes': [<matplotlib.lines.Line2D at 0x28d76d84e20>],
```

```
'medians': [<matplotlib.lines.Line2D at 0x28d72bbb760>],
'fliers': [<matplotlib.lines.Line2D at 0x28d72bbb160>],
'means': []}
```



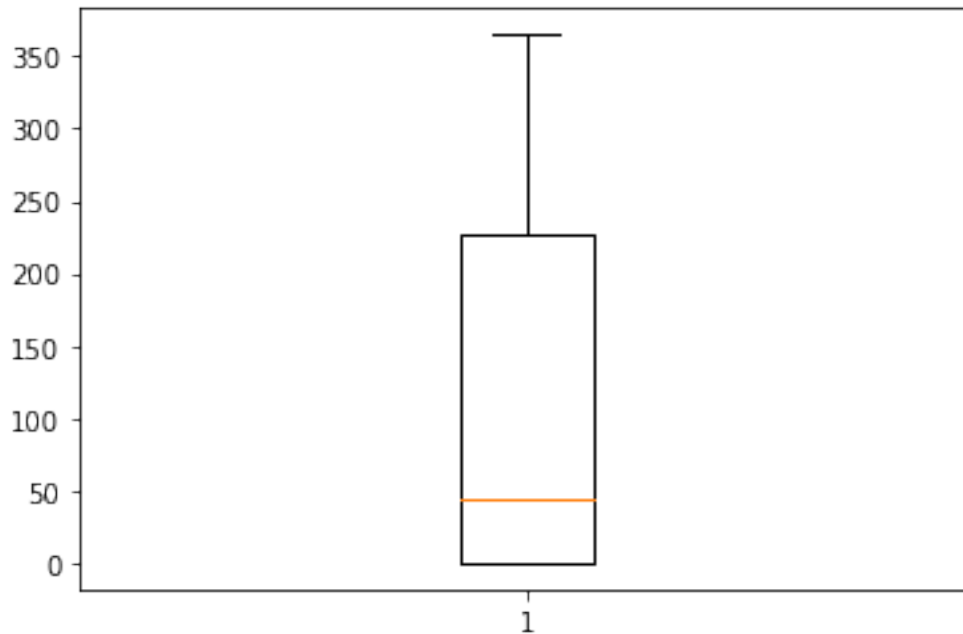
```
[115]: fig, ax = plt.subplots()
ax.boxplot(airbnb["minimum_nights"])
```

```
[115]: {'whiskers': [<matplotlib.lines.Line2D at 0x28d727cd250>,
<matplotlib.lines.Line2D at 0x28d727cd6a0>],
'caps': [<matplotlib.lines.Line2D at 0x28d727cd7f0>,
<matplotlib.lines.Line2D at 0x28d727cd790>],
'boxes': [<matplotlib.lines.Line2D at 0x28d72b42a30>],
'medians': [<matplotlib.lines.Line2D at 0x28d00005040>],
'fliers': [<matplotlib.lines.Line2D at 0x28d00005910>],
'means': []}
```



```
[116]: fig, ax = plt.subplots()
       ax.boxplot(airbnb["availability_365"])
```

```
[116]: {'whiskers': [<matplotlib.lines.Line2D at 0x28d7280e520>,
                  <matplotlib.lines.Line2D at 0x28d7280e130>],
       'caps': [<matplotlib.lines.Line2D at 0x28d027aaaf0>,
                <matplotlib.lines.Line2D at 0x28d027aa3d0>],
       'boxes': [<matplotlib.lines.Line2D at 0x28d7280ec10>],
       'medians': [<matplotlib.lines.Line2D at 0x28d027aa5b0>],
       'fliers': [<matplotlib.lines.Line2D at 0x28d027aa280>],
       'means': []}
```



- describe what you expected to see with these features and what you actually observed

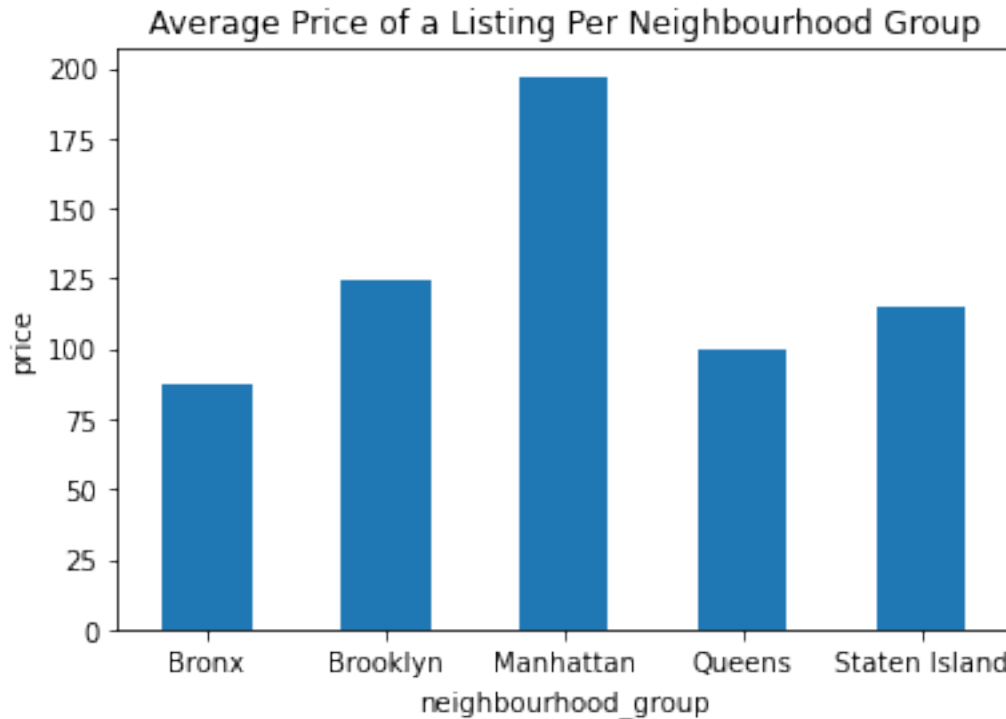
[Response here] * Price: I expected the prices to be skewed to the right since there are a lot of mansions that might cost a lot of money. * Minimum Nights: The number of minimum nights was also skewed to the right. * Availability 365: This field means the number of days a year the Airbnb is available for, and I expected that this value be less than or equal to 365 since there are only 365 days a year.

High variability in price with long tail values, review numbers much more compact, however availability has a wider variance.

3.0.3 [10 pts] Plot average price of a listing per neighbourhood_group

```
[117]: airbnb_gb_neighbourhood_group = airbnb.groupby(["neighbourhood_group"])
avg_by_neighbourhood_group = airbnb_gb_neighbourhood_group["price"].agg([np.
    ↳ average]).reset_index()
avg_by_neighbourhood_group.plot(kind="bar", x="neighbourhood_group",
    ↳ y="average", ylabel="price", rot=0, legend=False, title="Average Price of a
    ↳ Listing Per Neighbourhood Group")
```

```
[117]: <AxesSubplot:title={'center': 'Average Price of a Listing Per Neighbourhood
Group'}, xlabel='neighbourhood_group', ylabel='price'>
```



- describe what you expected to see with these features and what you actually observed

[Response here] As we can see from the bar chart above, Manhattan has the most expensive listings on average while the Bronx, in which 30.7% of the population lives below the poverty line, has the cheapest listings on average.

- So we can see different neighborhoods have dramatically different pricepoints, but how does the price breakdown by range. To see let's do a histogram of price by neighborhood to get a better sense of the distribution.

```
[99]: # see above
```

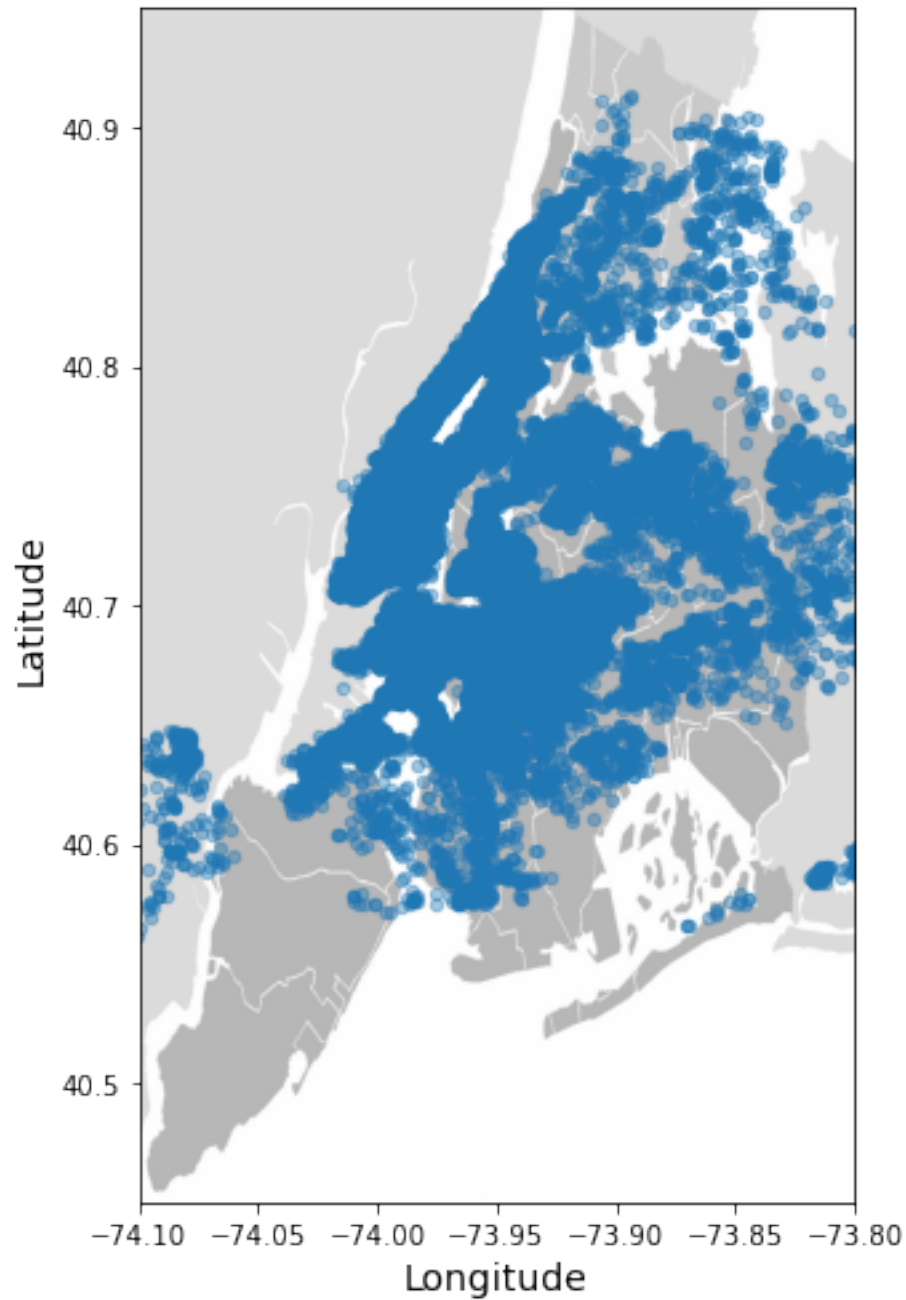
3.0.4 [5 pts] Plot map of airbnbs throughout New York (if it gets too crowded take a subset of the data, and try to make it look nice if you can :)).

```
[118]: images_path = os.path.join('./', "images")
os.makedirs(images_path, exist_ok=True)
filename = "newyork.png"

import matplotlib.image as mpimg
newyork_img=mpimg.imread(os.path.join(images_path, filename))
ax = airbnb.plot(kind="scatter", x="longitude", y="latitude", figsize=(10,7),
↳alpha=0.4)
```

```
# overlay the New York map on the plotted scatter plot  
# note: plt.imshow still refers to the most recent figure  
# that hasn't been plotted yet.  
plt.imshow(newyork_img, extent=[-74.10, -73.80, 40.45, 40.95], alpha=0.5)  
plt.ylabel("Latitude", fontsize=14)  
plt.xlabel("Longitude", fontsize=14)  
  
save_fig("newyork_airbnb_plot")  
plt.show()
```

Saving figure newyork_airbnb_plot



3.0.5 [5 pts] Plot average price of room types who have availability greater than 180 days and neighbourhood_group is Manhattan

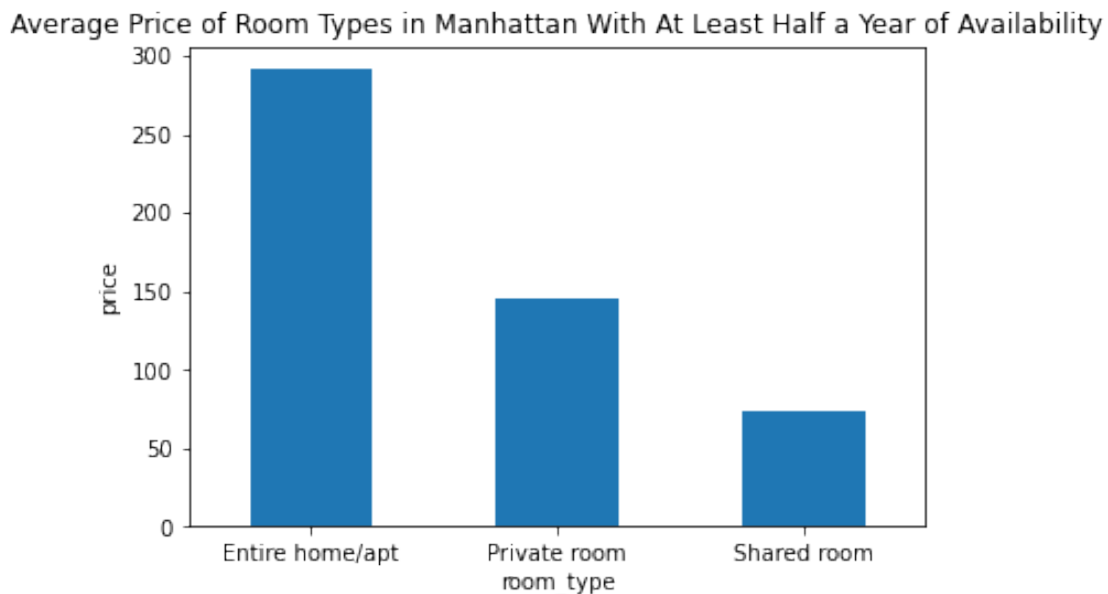
```
[119]: # bar chart
        # x-axis room types
        # y-axis average price
```

```

filtered_airbnb = airbnb[(airbnb["availability_365"] > 180) &
    ↪(airbnb["neighbourhood_group"] == "Manhattan")]
filtered_airbnb_gb_room_type = filtered_airbnb.groupby(["room_type"])
avg_by_room_type = filtered_airbnb_gb_room_type["price"].agg([np.average]).
    ↪reset_index()
avg_by_room_type.plot(kind="bar", x="room_type", y="average", ylabel="price",
    ↪rot=0, legend=False, title="Average Price of Room Types in Manhattan With At
    ↪Least Half a Year of Availability")

```

[119]: <AxesSubplot:title={'center': 'Average Price of Room Types in Manhattan With At Least Half a Year of Availability'}, xlabel='room_type', ylabel='price'>



3.0.6 [5 pts] Plot correlation matrix

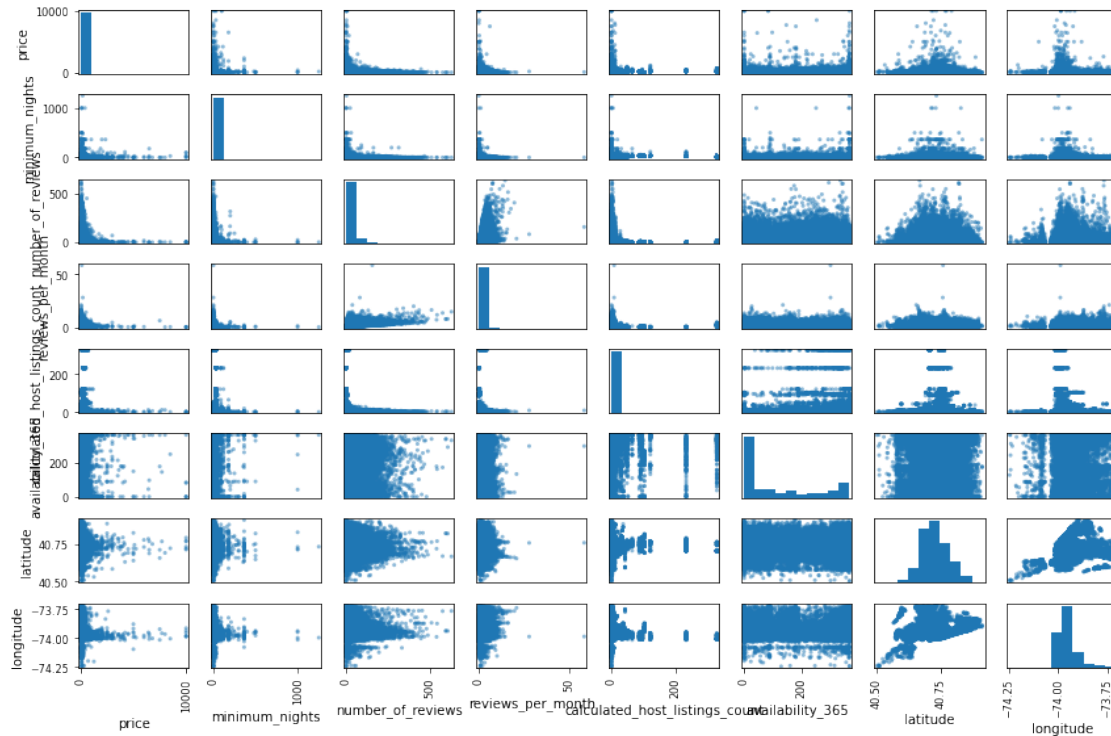
- which features have positive correlation?
- which features have negative correlation?

```

[120]: attributes = ["price", "minimum_nights", "number_of_reviews",
    ↪"reviews_per_month", "calculated_host_listings_count", "availability_365",
    ↪"latitude", "longitude"]
#airbnb[attributes].corr()
scatter_matrix(airbnb[attributes], figsize=(12, 8))
save_fig("airbnb_scatter_matrix_plot")

```

Saving figure airbnb_scatter_matrix_plot



[Response here] * number_of_reviews and reviews_per_month have a positive correlation as expected since they are related to one another mathematically. * It also seems like calculated_host_listings_count and availability_365 are positively correlated. A possible explanation for this is that if an airbnb is available more often in the year, the host might have to post each time they need to find a new customer. * availability_365 and reviews_per_month are also positively correlated since more availability means possibility for more customers. This same idea applies for availability_365 and number_of_reviews. * availability_365 and minimum_nights are positively correlated as well. A possible explanation for this is that if an Airbnb is available for larger portion of the year, the host will probably want to set the minimum_nights requirements to be higher so that they don't have to keep changing tenants all the time. * There is a negative correlation between longitude and price, implying that in general, airbnb's to the west (smaller longitude) are more expensive. * There are also more reviews_per_month for airbnb's to the east (greater longitude).

4 [30 pts] Prepare the Data

4.0.1 [5 pts] Augment the dataframe with two other features which you think would be useful

```
[121]: # There is a positive correlation between availability_per_listing and price.
# If availability_365 is small and the number of listings is large, then this
# ratio will be small.
```

```

# The price of the Airbnb will likely be small as well because if it's not
↳available for much of the year,
# and the host has to post multiple times, it probably means that there is less
↳interest in the Airbnb.
# It makes sense that the low demand might drive the prices down.
airbnb["availability_per_listing"] = airbnb["availability_365"] /
↳airbnb["calculated_host_listings_count"]

# There is a slight positive correlation between minimum_nights_per_listing and
↳price.
# A possible explanation for this is that if an Airbnb has a small number of
↳minimum nights, and the host
# has listed it multiple times, that probably means the Airbnb is struggling to
↳find tenants. Thus, the price
# will likely have to be dropped.
airbnb["minimum_nights_per_listing"] = airbnb["minimum_nights"] /
↳airbnb["calculated_host_listings_count"]

"""
airbnb["rpm_per_longitude"] = airbnb["reviews_per_month"] / airbnb["longitude"]

airbnb["lat_over_long"] = airbnb["latitude"] / airbnb["longitude"]
airbnb["lat_over_rpm"] = airbnb["latitude"] / airbnb["reviews_per_month"]
airbnb["long_over_availability"] = airbnb["longitude"] /
↳airbnb["availability_365"]
airbnb["lat_over_availability"] = airbnb["latitude"] / airbnb["availability_365"]
airbnb["rpm_per_listing"] = airbnb["reviews_per_month"] /
↳airbnb["calculated_host_listings_count"]
airbnb["rpm_per_min_night"] = airbnb["reviews_per_month"] /
↳airbnb["minimum_nights_per_listing"]
airbnb["num_reviews_over_rpm"] = airbnb["number_of_reviews"] /
↳airbnb["reviews_per_month"]
airbnb["availability_per_minimum_night"] = airbnb["availability_365"] /
↳airbnb["minimum_nights"]

airbnb["longitude_over_listings"] = airbnb["longitude"] /
↳airbnb["calculated_host_listings_count"]
airbnb["longitude_over_minimum_nights"] = airbnb["longitude"] /
↳airbnb["minimum_nights"]
airbnb["latitude_over_minimum_nights"] = airbnb["latitude"] /
↳airbnb["minimum_nights"]

airbnb["reviews_per_listing"] = airbnb["number_of_reviews"] /
↳airbnb["calculated_host_listings_count"]
airbnb["reviews_over_availability"] = airbnb["number_of_reviews"] /
↳airbnb["availability_365"]

```

```

airbnb["lat_over_listings"] = airbnb["latitude"] /
    ↳ airbnb["calculated_host_listings_count"]
"""
# obtain new correlations
corr_matrix = airbnb.corr()
corr_matrix["price"].sort_values(ascending=False)

```

```

[121]: price                1.000000
availability_per_listing    0.089530
availability_365            0.081829
calculated_host_listings_count 0.057472
minimum_nights              0.042799
minimum_nights_per_listing   0.040203
latitude                   0.033939
id                          0.010619
reviews_per_month           -0.030608
number_of_reviews           -0.047954
longitude                   -0.150019
Name: price, dtype: float64

```

4.0.2 [5 pts] Impute any missing feature with a method of your choice, and briefly discuss why you chose this imputation method

[]:

```

[134]: sample_incomplete_rows = airbnb[airbnb.isnull().any(axis=1)].head()

# I chose to remove any rows with a null value for "reviews_per_month" because
↳ that likely meant that number_of_reviews was 0,
# and it wouldn't make sense to fill in reviews_per_month with the median value
↳ when the total number of reviews that Airbnb actually
# got was 0.
sample_incomplete_rows.dropna(subset=["reviews_per_month"], inplace=True)
sample_incomplete_rows.head()

airbnb_features = airbnb.drop("price", axis=1) # drop labels for training set
↳ features
                                                    # the input to the model
↳ should not contain the true label
airbnb_labels = airbnb["price"].copy()

```

```

# I chose to remove any rows with a null value for "reviews_per_month" because
↳ that likely meant that number_of_reviews was 0,
# and it wouldn't make sense to fill in reviews_per_month with the median value
↳ when the total number of reviews that Airbnb actually
# got was 0.
airbnb.dropna(subset=["reviews_per_month"], inplace=True)
airbnb = airbnb[airbnb['calculated_host_listings_count'] != 0]

imputer = SimpleImputer(strategy="median") # use median imputation for missing
↳ values
categorical = ["neighbourhood", "neighbourhood_group", "room_type"]
airbnb_num = airbnb_features.drop(categorical, axis=1) # remove the categorical
↳ feature
airbnb_num.head()

```

```

[134]:
   id  latitude  longitude  minimum_nights  number_of_reviews  \
0  2539  40.64749  -73.97237             1             9
1  2595  40.75362  -73.98377             1            45
3  3831  40.68514  -73.95976             1           270
4  5022  40.79851  -73.94399            10             9
5  5099  40.74767  -73.97500             3            74

   reviews_per_month  calculated_host_listings_count  availability_365  \
0              0.21              6              365
1              0.38              2              355
3              4.64              1              194
4              0.10              1               0
5              0.59              1             129

   availability_per_listing  minimum_nights_per_listing
0              60.833333              0.166667
1             177.500000              0.500000
3             194.000000              1.000000
4               0.000000             10.000000
5             129.000000              3.000000

```

4.0.3 [15 pts] Code complete data pipeline using sklearn mixins

```

[141]: # column index
lat_idx, long_idx, min_nights_idx, num_reviews_idx, reviews_per_month_idx,
↳ listings_idx, availability_idx = 1, 2, 3, 4, 5, 6, 7

#
class AugmentFeatures(BaseEstimator, TransformerMixin):
    '''
    implements the previous features we had defined

```

```

        housing["rooms_per_household"] = housing["total_rooms"] /
        ↪ housing["households"]
        housing["bedrooms_per_room"] = housing["total_bedrooms"] /
        ↪ housing["total_rooms"]
        housing["population_per_household"] = housing["population"] /
        ↪ housing["households"]
        '''
    def __init__(self, add_min_nights_per_listing = True):
        self.add_min_nights_per_listing = add_min_nights_per_listing

    def fit(self, X, y=None):
        return self # nothing else to do

    def transform(self, X):
        availability_per_listing = X[:, availability_idx] / X[:, listings_idx]
        if self.add_min_nights_per_listing:
            min_nights_per_listing = X[:, min_nights_idx] / X[:, listings_idx]
            return np.c_[X, availability_per_listing, min_nights_per_listing]
        else:
            return np.c_[X, availability_per_listing]

#attr_adder = AugmentFeatures()
#airbnb_extra_attribs = attr_adder.transform(airbnb.values) # generate new
↪ features

# this will be a numerical pipeline
# 1. impute, 2. augment the feature set 3. normalize using StandardScaler()
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attribs_adder', AugmentFeatures()),
    ('std_scaler', StandardScaler()),
])

#housing_num_tr = num_pipeline.fit_transform(housing_num)

#numerical_features = [x for x in airbnb_num.columns.tolist() if x not in
↪ categorical]
numerical_features = airbnb_num.columns.tolist()
categorical_features = categorical

full_pipeline = ColumnTransformer([
    ("num", num_pipeline, numerical_features),
    ("cat", OneHotEncoder(), categorical_features),
])

airbnb_prepared = full_pipeline.fit_transform(airbnb)
airbnb_prepared

```

```
[141]: <38843x238 sparse matrix of type '<class 'numpy.float64'>'
      with 582645 stored elements in Compressed Sparse Row format>
```

4.0.4 [5 pts] Set aside 20% of the data as test test (80% train, 20% test).

```
[142]: data_target = airbnb['price']
      train, test, target, test_target = train_test_split(airbnb_prepared,
      ↪data_target, test_size=0.2, random_state=0)
```

5 [15 pts] Fit a model of your choice

The task is to predict the price, you could refer to the housing example on how to train and evaluate your model using MSE. Provide both test and train set MSE values.

```
[143]: lin_reg = LinearRegression()
      lin_reg.fit(train, target)

      # let's try the full preprocessing pipeline on a few training instances
      data = test
      labels = target_test

      print("Predictions:", lin_reg.predict(data)[:5])
      print("Actual labels:", list(labels)[:5])
```

```
Predictions: [216.63193311 291.71321556  86.98501491  67.89549236 153.10099963]
Actual labels: [136900.0, 241300.0, 200700.0, 72500.0, 460000.0]
```

```
[145]: preds = lin_reg.predict(test)
      mse = mean_squared_error(test_target, preds)
      rmse = np.sqrt(mse)
      rmse
```

```
[145]: 163.3978324537358
```